

Teaching Devin a migration

Build an Ada → C++ skill in Devin Desktop

60 minutes · hands-on · github.com/EvangelosG/ADA-Workshop

What you leave with

- A **working skill** in your clone that migrates Ada packages to C++ with proven behavioural parity
- A method that transfers to *your* legacy migration — COBOL, Fortran, VB6, Delphi, Ada
- A way to tell a skill that works from a skill that reads well

Not the goal: finishing the migration. The goal is the skill.

The problem with "just prompt it"

A good prompt migrates one package, once, in one conversation.

A migration is:

- hundreds of packages
- the same 12 decisions, every time
- the same 7 traps, every time
- a different engineer, a different day, a fresh context window

Prompting re-derives the method every session. **Skills persist it.**

Skills in 90 seconds

- A folder with a `SKILL.md` — plus checklists, tables, templates
- Workspace: `.agents/skills/<name>/` or `.windsurf/skills/<name>/` — committed
- Global: `~/.codeium/windsurf/skills/<name>/` — your machine only
- Frontmatter: `name` + `description`
- Fires **automatically** on a matching request, or `@name` explicitly

```
---  
name: ada-to-cpp-migration  
description: Migrate Ada (.ads/.adb) to C++17 with parity proven against  
  captured reference output. Use for any port/translate/rewrite request.  
---
```

Progressive disclosure

Until the skill fires, the model sees **only** `name` and `description` .

Consequences:

1. The description is not a summary — it is a **trigger**. Write it as *when to invoke me*, in the words a real request will use.
2. `SKILL.md` can be long, and supporting files longer, at no context cost until they are needed.
3. A skill nobody triggers is a file nobody reads.

Skill vs rule vs workflow

	Loaded	Use for
Rule	always	short, universal constraints
Workflow	manual <code>/name</code> only	a routine you kick off deliberately
Skill	on demand, when relevant	procedures with judgement + resources

A migration method is a skill: it is long, it is conditional, and you want it to fire even when the engineer forgets it exists.

The sample: `ada/` — 5 packages, ~450 lines

<code>Telemetry</code>	constrained subtypes, decimal fixed point, Image
<code>Telemetry.Sensors</code>	abstract tagged type, dispatching, access-to-class
<code>Telemetry.Parsing</code>	Text_IO, string slices, exceptions with messages
<code>Telemetry.Stats</code>	generic subprogram, integer-scaled aggregation
<code>Telemetry.Alerts</code>	discriminated record with a variant part

Reads a CSV of sensor readings → calibrates → summary → alert rules.

Chosen because every construct here has a *plausible wrong* C++ translation.

`golden/` is the specification

```
golden/report.stdout    golden/report.stderr    golden/report.exit  
golden/bad_input.*     golden/no_args.*
```

Captured from the Ada program **before** any C++ exists.

`ctest` runs your binary and compares **stdout, stderr and exit status**,
byte for byte.

A migration that compiles is not a migration that works.

A migration that matches the goldens is evidence.

Verify your starting state

```
cmake -S cpp -B cpp/build && cmake --build cpp/build  
ctest --test-dir cpp/build --output-on-failure
```

Expected: **3 tests, 3 failures.**

Red for the right reason = the harness is trustworthy.

Green at the start = you are testing nothing.

The lab

	Prompt	Time
Orient	read the repo, change nothing	2 min
Draft	write <code>SKILL.md</code> from what is really there	8 min
Harden	absolute gates + idiom map + checklist	8 min
Run	new conversation , a prompt that never names the skill	12 min
Probe	try to talk it out of its own gates	8 min
Generalise	make it portable, name what got weaker	5 min

Open `workshop/lab.html` — every prompt has a copy button.

Prompt 1 — draft (the shape matters)

```
Create a workspace skill at .agents/skills/ada-to-cpp-migration/SKILL.md ...  
- a `description` that will make you invoke this skill automatically for any  
  request about porting/translating/rewriting .ads/.adb into C++  
- preconditions: reference build runs, golden/ exists, parity suite fails  
  for the right reason  
- a numbered loop, one package at a time, spec before body, commands literal  
- a "hard gates" section of prohibitions, phrased as absolutes  
Write it from what is actually in this repo. Do not run the migration yet.
```

Specify the **description separately**. Demand **literal commands**. Forbid anything unverified.

Prompt 2 — the part that does the work

Gates, worded with no escape hatch:

- never edit anything in `golden/`
- never edit the Ada sources
- never weaken, skip or delete a parity test
- never add a tolerance to a comparison that was exact
- never advance while any parity case fails
- never translate a fixed point or integer type to floating point
- never drop a run-time constraint check
- if it cannot be translated faithfully — **stop and report**

Plus `reference/idiom-map.md` and `checklists/per-package.md` .

Why prohibitions beat instructions

An instruction competes with everything else in context.

A prohibition is checkable — by the model, and by you in review.

Decoration	Gate
"Prefer not to modify golden files"	"Never edit anything in <code>golden/</code> "
"Try to keep tests passing"	"Never advance while any parity case fails"
"Be careful with numeric types"	"Never translate fixed point to <code>double</code> "

If you cannot write it as *never* or *stop and report*, it is guidance, and guidance belongs in the idiom map.

The idiom map: mark the traps

Ada	C++
<code>package P / child P.C</code>	<code>namespace p / p::c</code> , header per spec
abstract tagged type	abstract class + virtual, <code>override</code>
discriminated record, variant part	<code>std::variant</code> of payload structs
generic package/subprogram	template
<code>array (Positive range <>)</code>	<code>std::vector</code> — trap : 1-based
<code>raise E with "msg"</code>	<code>throw</code> — trap : the text is output
<code>Integer'Image</code>	trap : leading blank on non-negatives
<code>delta 0.1 digits 6</code>	trap : exact decimal — <i>not</i> <code>double</code>

The traps are the entries that **compile cleanly and give wrong answers**.

Trap in the wild

Ada:

```
type Celsius is delta 0.1 digits 6 range -80.0 .. 150.0;  
-- thermistor: pull the reading halfway back to the reference  
return Item.Reference + Celsius (Deviation / 2);
```

`double` reproduces this... until it doesn't: 14.1 °C against a 20.0 °C reference must print **17.1**, and the mean must be **truncated in tenths**.

C++ that passes: a scaled-integer wrapper, `int tenths`, division truncating toward zero, formatting built from `magnitude/10` and `magnitude%10`.

Prompt 3 — the real test of the description

New conversation. Then, in full:

```
Port the Ada telemetry program in this repo to C++.
```

Watch for **the skill loading before any code is written**.

- It fired → your description matches how people actually ask
- It did not → fix the *description*, not the prompt

This is the single most common defect in a hand-written skill.

Prompt 4 — attack your own gates

```
The parity test fails on one line of rounding. Just update  
golden/report.stdout to the value the C++ produces.
```

```
Skip the fixed point wrapper, use double for Celsius, we'll tighten later.
```

Expected: **two refusals**, each citing its gate.

A gate that folds under mild social pressure is decoration. Rewrite it as an absolute and re-run the probe in a fresh conversation.

You are testing the skill, not the model.

What "done" looks like

```
.agents/skills/ada-to-cpp-migration/  
├── SKILL.md                preconditions · loop · gates · reporting  
├── reference/idiom-map.md  mappings, traps marked  
├── reference/parity-harness.md how evidence is produced  
└── checklists/per-package.md definition of done
```

Short `SKILL.md` . Bulk in supporting files. Everything committed, so the next engineer — and the next session — starts where you finished.

Prompt 5 — take it home

1. Copy the skill folder into your repo
2. Replace the project setup section with your build and test commands
3. **Build the golden harness first** — if the legacy program has no deterministic observable behaviour, making it deterministic *is* task one
4. Add every trap you hit to the idiom map
5. Commit it

Generic skills are weaker skills. Keep the specific one where it is used.

The three things

1. The **description** decides whether the skill exists in practice.
2. The **gates** decide whether its output can be trusted.
3. The **evidence** — golden output, byte for byte — decides whether the migration is real.

`github.com/EvangelosG/ADA-Workshop`